

# Parallel Corner Detection



Kunal Jobanputra and Ryan Stentz

# Applications of the Harris Corner Detector

Optical  
Flow



Rotating camera, arbitrary world



Homography  
Estimation



Object  
Recognition



Feature  
Matching



# Harris Corner Detection Algorithm

Original



Gaussian Blur + Partial




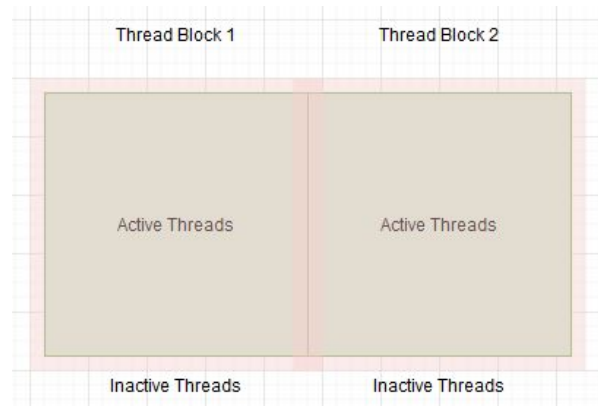
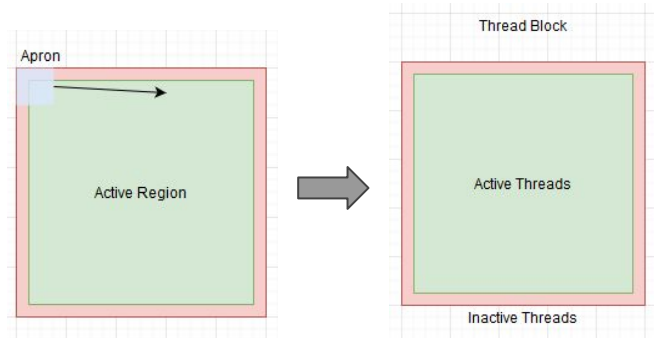
Corners Detected



# GPU - Approaches

## Partial Derivatives (Sobel Convolution):

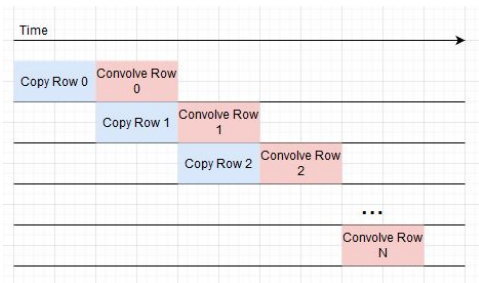
- Unroll Kernel Loop
  - Reduces data accesses and arithmetic
- Utilize shared memory
  - Each thread loads their pixel into memory
  - Creates idle threads
  - More thread blocks per image
- Utilize shared memory - 2nd approach
  - Each thread loads their pixel into memory, some threads load an additional pixel in the apron
  - No idle threads



# GPU - Approaches

## Partial Derivatives (Sobel Convolution): Streaming approach

- Helps hide memory overhead
- Utilizes CUDA streams to copy and process rows of an image independently



## Cornerness Computation:

- Utilized the shared memory approach on the previous slide
- Compute and use integral images (pictured below)

1.

31	2	4	33	5	36
12	26	9	10	29	25
13	17	21	22	20	18
24	23	15	16	14	19
30	8	28	27	11	7
1	35	34	3	32	6

2.

31	33	37	70	75	111
43	71	84	127	161	222
56	101	135	200	254	333
80	148	197	278	346	444
110	186	263	371	450	555
111	222	333	444	555	666

$$15 + 16 + 14 + 28 + 27 + 11 =$$

$$101 + 450 - 254 - 186 = 111$$

## Output Compression:


- Switch from a float image to a byte image
- Compress the output into a vector of pixel locations instead of a binary image mask

# GPU - Results

Our most optimized GPU implementation consisted of:

1. Shared memory sobel filter
2. Shared memory cornerness calculator
3. Non-maximal suppression
4. Slightly compressed output

We achieved a speedup of more than 3x, making 1080p and 4k image processing on embedded devices feasible in real time



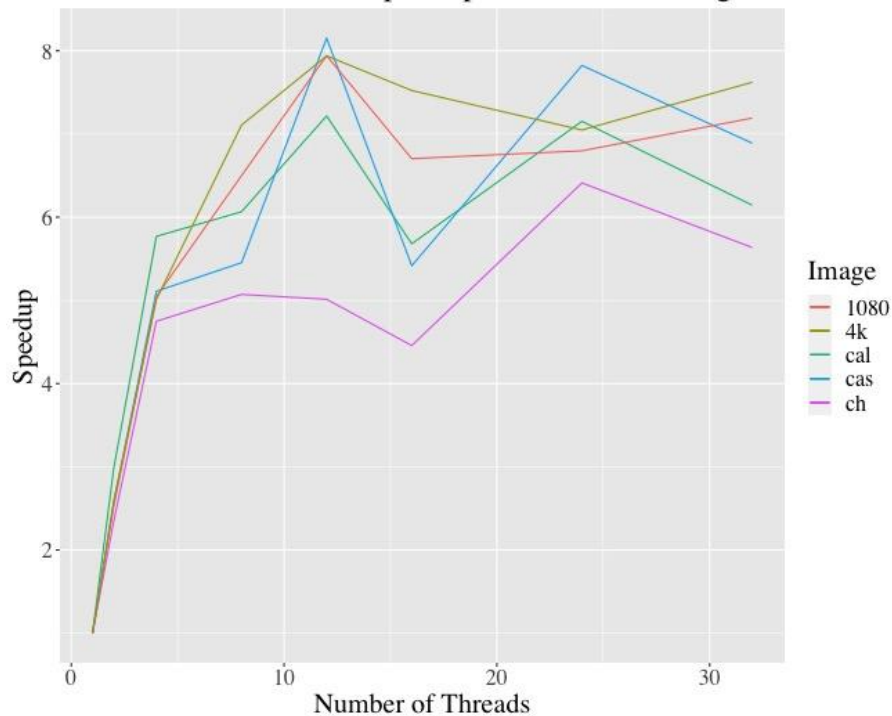
# Shared Memory - Approach

- Static assignment
- Rewriting the loop to better use resources

```
#pragma omp parallel for num_threads(numThreads)
for (int it = 0; it < srcGray.rows * srcGray.cols; it++) {
    int i = it / srcGray.cols;
    int j = it % srcGray.cols;
    harris.at<float>(i0: i, i1: j) = c(i, j);
}
```

# Shared Memory: Up to 8.15x Speedup

Number of Threads vs Speedup for Different Images



- Initial superscalar speedup
- Highest speedup around 12 threads, with one anomaly
- Spikes again round 24 threads, with anomalies from bigger images
- A tough apples to apples comparison to GPU implementation due to cost disparity

